

ECE 340 Lab 4

Omar Mahmoud

1753607

Section D21

11/18/2024

1 Low-Pass Filtering Audio Signals

The script below creates a 513-tap low-pass FIR filter with a cut-off frequency of 2500 Hz and a sampling frequency of 22050 Hz. The truncation window was selected to ensure that the stop-band ripples of the frequency response did not exceed -50dB:

```
figure(1);
fc = 2500;           % Cutoff frequency in Hz
Fs = 22050;          % Sampling frequency in Hz
wc = fc / (Fs / 2);  % Normalized cutoff frequency

% Windowing parameters
window = hamming(513); % Hamming window of length 513

% Design the FIR filter using the fir1 function
filter_coeff = fir1(513 - 1, wc, 'low', window); % 513 taps, 513 - 1 in the
order

% Plot the frequency response of the filter
freqz(filter_coeff, 1); % 1024 points for a smooth plot, Fs as sampling
frequency
```

The code above generates the following frequency response plots:




Figure 1: Plot of the Frequency Response of the Low-Pass FIR filter.

Given the audio file *love_mono22.wav* the script below reads the audio file, passes it through the FIR filter, and calculates the output signal:

```
figure(2);
% Read audio file
[x, Fs] = audioread('love_mono22.wav');

% Filter the audio signal
x_filtered = filter(filter_coeff, 1, x);

% Plot the inputs power spectral density
subplot(2, 1, 1);
pwelch(x, [], [], [], Fs);
title('Power Spectrum of Input Signal');
```

```

% Plot the outputs power spectral density
subplot(2, 1, 2);
pwelch(x_filtered, [], [], [], Fs);
title('Power Spectrum of Filtered Signal');

% Create output file
audiowrite('filtered_love_mono22.wav', x_filtered, Fs);

```

The script above produced the following power spectral density plots:




Figure 2: Power Spectrums of Input (top) and Output (bottom) Signals.

Listening to the original and filtered signals it is clear that the filtered signal is now bass heavy in comparison to the input signal. This is because the low-pass filter attenuates high frequency signals while passing low frequency signals attributed to bass. It's also clear that the constant tone that can be heard throughout the song was filtered significantly as it's much less noticeable in the output signal. This is supported by analyzing the power spectrums of both signals. While there is a clear peak at approximately 3 kHz in the input signal, this peak is drastically reduced in the output signal.

2 High-Pass Filtering Audio Signals

The script below creates a 513-tap high-pass FIR filter with a cut-off frequency of 5000 Hz and a sampling frequency 5000 22050 Hz. The truncation window was selected to ensure that the stop-band ripples of the frequency response did not exceed -50dB. It then once again reads the *love_mono22.wav* audio file, passing its signal into the high-pass filter, and plotting the power density spectrums. Finally the new signal is outputted to a new *filtered_love_mono22.wav* file:

```

figure(1);
fc = 5000;           % Cutoff frequency in Hz
Fs = 22050;          % Sampling frequency in Hz
wc = fc / (Fs / 2);  % Normalized cutoff frequency

% Windowing parameters
window = hamming(513); % Hamming window of length 513

% Design the FIR filter using the fir1 function

```

```

filter_coeff = fir1(513 - 1, wc, 'high', window); % 513 taps, 513 - 1 in the
order

% Plot the frequency response of the filter
freqz(filter_coeff, 1); % 1024 points for a smooth plot, Fs as sampling
frequency

figure(2);
% Read audio file
[x, Fs] = audioread('love_mono22.wav');

% Filter the audio signal
x_filtered = filter(filter_coeff, 1, x);

% Plot the inputs power spectral density
subplot(2, 1, 1);
pwelch(x, [], [], [], Fs);
title('Power Spectrum of Input Signal');

% Plot the outputs power spectral density
subplot(2, 1, 2);
pwelch(x_filtered, [], [], [], Fs);
title('Power Spectrum of Filtered Signal');

% Create output file
audiowrite('filtered_love_mono22.wav', x_filtered, Fs);

```

The code above generates the following frequency response plots:




Figure 3: Plot of the Frequency Response of the High-Pass FIR filter.

And the following power density spectrums:




Figure 4: Power Spectrums of the Input (top) and Output (bottom) Signals.

The filtered audio signal is almost unrecognizable to the input signal. There vocals and bass have been nearly completely attenuated. The only recognizable elements from the original signal are the high frequency elements of the instrumentals. Similar to the low-pass filter, the constant tone that can be heard throughout the song was filtered significantly as it's much less noticeable in the output signal. This is supported by analyzing the power spectrums of both signals. While there is a clear peak at approximately 3 kHz in the input signal, this peak is drastically reduced in the output signal.

3 Audio Signal Filter Design

In order to filter out the constant tone without negatively affecting other aspects of the audio signal, it was passed through a band-stop filter per the script below:

```
figure(1);
fc_lower = 2900; % Lower cutoff frequency in Hz
fc_upper = 3100; % Upper cutoff frequency in Hz
Fs = 22050; % Sampling frequency in Hz
wc_lower = fc_lower / (Fs / 2); % Normalized lower cutoff frequency
wc_upper = fc_upper / (Fs / 2); % Normalized upper cutoff frequency

% Windowing parameters
window = hamming(513); % Hamming window of length 513

% Design the FIR filter using the fir1 function
filter_coeff = fir1(513 - 1, [wc_lower wc_upper], 'stop', window); % 513 taps,
513 - 1 in the order

% Plot the frequency response of the filter
freqz(filter_coeff, 1); % 1024 points for a smooth plot, Fs as sampling
frequency

figure(2);
% Read audio file
[x, Fs] = audioread('love_mono22.wav');

% Filter the audio signal
```

```

x_filtered = filter(filter_coeff, 1, x);

% Plot the inputs power spectral density
subplot(2, 1, 1);
pwelch(x, [], [], [], Fs);
title('Power Spectrum of Input Signal');

% Plot the outputs power spectral density
subplot(2, 1, 2);
pwelch(x_filtered, [], [], [], Fs);
title('Power Spectrum of Filtered Signal');

% Create output file
audiowrite('filtered_love_mono22.wav', x_filtered, Fs);

```

The code above generates the following frequency response plots:




Figure 5: Plot of the Frequency Response of the Band-Stop FIR filter.

And the following power density spectrums:




Figure 6: Plot of the Power Spectrums of the Input (top) and Output (bottom) Signals.

Listening to the output signal, the tone is substantially attenuated while there are no audible effects on the rest of the recording. This is also visible on the power density spectrum of the output where it is clear that the peak at 3 kHz is attenuated significantly.

4 Filtering Noisy Images

Given the image file *ayantika.tif*:




Figure 7: Original *ayantika.tif* Image.

The code below filters the image and plots the power spectrums of the input and output images:

```
% MATLAB code for spectral analysis and lowpass filtering of an image
% see section 17.6, Fig. 17.20, Fig. 17.21
%
% reading the image 'ayantika.tif'
I = imread('ayantika.tif');
```

```

Iu=I;          % copy of original image
imshow(I)      % display the image
fprintf('\nThe image Ayantika has been displayed')
fprintf('\nPress any key to continue')
pause
I = double(I);
I = I - mean(mean(I));
% 2D Bartlett window
x = bartlett(32);
for i = 1:32
    zx(i,:) = x';
    zy(:,i) = x ;
end
bartlett2D = zx .* zy;
%
n = 0;
% calculate power spectrum
P = zeros(256,256);
for (i = 1:16:320)
    for (j = 1:16:288)
        Isub = I(i:i+31,j:j+31).*bartlett2D;
        P = P + fftshift(fft2(Isub,256,256));
        n = n + 1;
    end
end
Pabs = (abs(P)/n).^2;
mesh([-128:127]*2/256, [-128:127]*2/256, Pabs/max(max(Pabs)));
xlabel('Horizontal Frequency'); ylabel('Vertical Frequency');
zlabel('Image Power Spectrum (in dB)');
print -dtiff plot1.tiff
fprintf('\nThe Image Power Spectrum has been displayed and saved')
fprintf('\nPress any key to continue')
pause

filter_coeff = [1  2 3 2 1; 2  3 4 3 2; 3  4 5 4 3; 2  3 4 3 2; 1  2 3
                2 1]/65 ;
% Frequency Response plot
spec = fft2(filter_coeff,128,128) ;      % Frequency spectrum of 2-D Filter
R = abs(spec(1:65,1:65)) ;
mesh(R), grid
xlabel('Horizontal Frequency')
ylabel('Vertical Frequency');
zlabel('Frequency Response');
print -dtiff plot.tiff
fprintf('\nThe Frequency Response of the 2-D Filter has been displayed and saved
')
fprintf('\nPress any key to continue')
pause
%
% Digital filtering the image
filtered_image = filter2(filter_coeff, double(I)) ;
imshow(uint8(filtered_image))
%imshow(Iu)
imwrite(uint8(filtered_image),'ayantika_filt.tif','tif') ;
fprintf('\nThe Filtered Image has been displayed and saved')

```

The script above returns the following plots and images:




Figure 8: Plot of the Power Spectrum of the Input Image.




Figure 9: Plot of the Power Spectrum of the Output Image.



Figure 10: *ayantika.tif* After Filtering.

The input image had a very prominent grid like artifact which separated the shaded pixels apart. The power spectrum showed the peaks/nodes that created the artifact. Applying a low-pass filter resulted in an image where lower frequency darker shades are more prominent, and all higher frequency lighter shades are attenuated.